

Design proposal for best effort pod autoscaling in kubernetes with less disruptions

1. Use node affinity and anti-affinity for scheduling of nodes.

Example:

1. case1: If we want to schedule a pod on nodes which has more resources, we can use operator 'In' in that case.
2. case2: If we do not want to schedule a pod on critical nodes which has no resources, we can use operator 'Not In' in this case.

In second case we can many types of label one with high, moderate and low.

Default Case:

- A process to update the label or value we use in affinity and anti-affinity
Note: We will not schedule the pods once it reaches 80% as per our calculation (RAM * 0.9) / 300 MB. 300MB can be changed.
- We can use settings **requiredduringschedulingignoreduringexecution**
- In future we can have our custom scheduler.

2. To be extra safe we can even taint the nodes which has more resource usage [it is redundant as we are updating labels already]

3. Use priority efficiently with pods

- Documentation says that which has more priority has a less chance of getting evicted.
- We can have a process which updates the pods priority based on resource utilization and other factors in a node.
- I didn't get a solution to update the pods priority while pod running.
- If update priority does not work, we can try updating of priority classes value itself.

Important Note: All this are just a solution based on research but must be tested. if it works can be improved or can be evicted [most likely rebalancer won't work because updating of spec is not supported].

Important Resource

1. Affinity / Anti-affinity of node
2. Taint
3. Priority Classes

related links

Priority Classes

<https://kubernetes.io/docs/concepts/configuration/pod-priority-preemption>

kubernetes best practices

<https://cloud.google.com/blog/products/gcp/kubernetes-best-practices-resource-requests-and-limits>

Affinity and anti-affinity and taint

<https://kubernetes.io/blog/2017/03/advanced-scheduling-in-kubernetes/>